Mohan Sri Raghuram

+91 9392380182 | msriraghuram@gmail.com | linkedin.com/in/msriraghuram | github.com/raghurammrsd

Education

Parul University — Vadodara, India
B.Tech in Computer Science and Engineering (AI) | CGPA: 8.21/10 — *Jul 2023 – Apr 2027*

Experience

Frontier Atlas – Founding Member | AI Engineer — Sep 2026 – Present
Products: frontieratlas.co | aiorbit.club — *Remote*

- Contribute to AI product development for Frontier Atlas and AIOrbit.club, working across backend systems and AI data pipelines using Python and PostgreSQL.
- Build LLM applications using RAG and NLP, and work with knowledge graphs to structure and connect data.
- Develop web data extraction and data quality workflows that support AI-driven products.

One Atlas – AI Engineer Intern — May 2026 – Present
Remote

- Built scalable Python data pipelines to ingest, validate, and structure technical data using AsyncIO, Crawl4AI, and PostgreSQL.
- Designed LLM-assisted information extraction workflows using Gemini and Groq to automate entity extraction, knowledge generation, and structured data enrichment.
- Developed evaluation and data quality pipelines for AI datasets, implementing schema validation, metadata extraction, and automated quality assessment.
- Designed reusable Python modules and backend components to support maintainable, scalable AI workflows.

Projects

CodeGraph MCP – Repository Intelligence Engine | *Python, MCP, SQLite, AST* | GitHub | PyPI

- Engineered a local-first MCP server that lets AI coding agents deterministically query codebases for symbols, dependencies, call graphs, routes, tests, Git relationships, database models, and runtime behavior instead of repeatedly grepping and reading entire files.
- Built a deterministic code-intelligence pipeline on Python AST/static analysis with an incremental SQLite index, resolving symbols, imports, callers/callees, API routes, and cross-file references with precise `file:line` evidence.
- Implemented evidence-backed retrieval and deterministic ranking that separates verified structural relationships from inferred or ambiguous ones, reducing hallucinated codebase understanding.
- Added database and runtime intelligence linking routes/services to ORM models, SQL queries, tables, columns, and observed runtime behavior, so agents can reason about both static architecture and actual execution paths.
- Optimized for large repositories with incremental indexing, bounded graph traversal, caching, and context-efficient structured outputs, cutting unnecessary file reads and agent tool calls.
- Validated via automated tests, static analysis, type checking, and repository-scale benchmarks on production-sized codebases; published on PyPI for use with MCP-compatible AI coding environments.

AI App Compiler Platform | *Python, FastAPI, Pydantic* | Live Demo | GitHub — May 2026 – Jun 2026

- Built an AI-powered application compiler that converts natural language requirements into software architectures through a deterministic multi-stage pipeline.
- Implemented Pydantic-based cross-layer validation with automated consistency checks across APIs, databases, RBAC, and workflows to catch integration errors early.
- Developed a FastAPI workbench with runtime simulation, CRUD validation, authentication testing, and architecture visualization.

DUNDOO – AI Hyperlocal Commerce Platform | *Flask, PostgreSQL* | Live Demo | GitHub — Nov 2025 – Feb 2026

- Architected a Flask and PostgreSQL backend for a hyperlocal commerce platform with map-based navigation and payment integration.
- Designed RESTful services for product search, inventory management, multilingual voice search, and payment workflows.

Technical Skills

Languages: Python, SQL, C
AI/ML: Machine Learning, Deep Learning, NLP, Computer Vision, CNN, LLMs, RAG, Prompt Engineering, MCP
Backend: Flask, FastAPI, REST APIs, AsyncIO
Databases: PostgreSQL, MySQL, MongoDB, SQLite, Vector Databases
Libraries & Tools: PyTorch, TensorFlow, OpenCV, Scikit-learn, Pandas, NumPy, Crawl4AI, Git, GitHub, Postman, Jupyter Notebook
Cloud & DevOps: AWS, Docker, Cloudflare Workers
Concepts: ETL Pipelines, Entity Resolution, Web Scraping, API Integration, Data Validation, Static Analysis

Achievements & Certifications

Achievements: Google Cloud Skill Badge – Machine Learning Engineer | Finalist, PU CODE Hackathon 3.0
Certifications: IBM AI Practitioner | Deloitte Data Analytics Simulation (Forage) | Data Science & Analytics (HP LIFE) | Web API Development with Flask (Infosys Springboard)